



Distributed Systems

March 5, 2001



Distributed Systems

What is a Distributed System?

- Tanenbaum and van Renesse: A distributed system is one that looks to its users like an ordinary, centralized, system but runs on multiple independent CPUs
- Symptoms? Shroeder:
 - Multiple, independent processing units
 - Processors communicate via a hardware interconnect
 - Processing unit failures are independent
 - Manage resource sharing
 - State is shared among processors

2



Distributed Systems, cont.

Design Issues

- | | |
|-----------------|-----------------|
| ▪ Scaling | ▪ Consistency |
| ▪ Communication | ▪ Failures |
| ▪ Coordination | ▪ Security |
| ▪ Transparency | ▪ Heterogeneity |
| ▪ Naming | ▪ Mobility |
| ▪ Load sharing | |

3



Distributed Systems, cont.

Naming

- What do we need to name?
 - Processes, Services, Hosts, Objects, Groups, ...
- Binding - Discovering and associating a name to a one of these
- e.g., DNS: hostname to IP address

4



Communication

- Messages can have many characteristics:
 - Length, priority, streams
- Communication medium properties affect communication performance
 - bandwidth, latency, multi-cast capability, message prioritization

5



Consistency

- Since we assume network links can fail at any time, replication is required to maintain consistency for longer computations
 - Replication of data
 - Replication of computation
- Costs associated with consistency:
 - Reduction in the amount of effective resources
 - Managing extended failures

6



Load Sharing

- Local vs. non-local
 - e.g., communication failure less likely in local clusters of processors
- Process migration can be expensive
- What about a system that knows the load of each machine, then assigns computation?
 - Doesn't scale well
 - Issues - Turning Completeness, Propagation Delay

7



Remote Procedure Call

- Introduced by Birrell and Nelson in 1985 (Xerox PARC)
- Abstracts the notion of where a computation runs
- Semantics can include:
 - notion that computation may fail (e.g., errors, faults, deadlocks)
 - blocking or non-blocking
 - synchronous vs. asynchronous
 - at least once, at most once, exactly once

8



Remote Procedure Call, cont.

- Implementation assumes a client/server framework
- RPC infrastructure must implement:
 - Server Registration, Binding, Marshalling, Message Send/Receive
 - Compiler/linker can implement stubs for RPC
- RPC used by NFS, AFS, but not X

9



Server Registration

- Use Discovery or Directory Service
- Each server registers a service and a version
 - Service - typically implemented as an integer; a set of well-known services exist
 - Version - a number used to select which server to use
- Who decides which server to return to the client?
 - Key issue: Load Balancing

10



Marshalling

- Assume heterogeneous computing environment
- Need to have a consistent set of data types
- What about pointers?
 - Client/Server operate in separate contexts...
- What about exceptions?
 - Exceptions can result from errors at the client site or network problems

11



End to End Design

- History: Design of the Internet
 - Q: Where should functionality be built into the Internet
 - e.g., should sessions & flows be a fundamental part of the network, or something implemented at higher layers?

12



End to End Design

- Insight - Many functions can be built at the ends of the network (e.g., RPC). Thus, building these functions in at lower layers will have significant tradeoffs
 - e.g., In the 70's, everyone used the DARPA-net to log into computers. If "sessions" had been built into the lower layers of the Internet, web-based traffic in the 90's would not have scaled well